

# Mixed-Precision Inference on GPUs

## Benchmarking FP32, FP16, and BLAS-Optimized Neural Network Inference

### GPU Computing Project

Option 2 — Mixed Precision with Row-Wise Scaling

February 27, 2026

# Outline

- 1 Problem & Objectives
- 2 Background Concepts
- 3 The Six Precision Modes
- 4 Architecture & Implementation
- 5 Key Algorithms
- 6 Metrics & Methodology
- 7 Optimizations & Challenges
- 8 Results & Analysis
- 9 Conclusions & Future Work

# Problem Statement

## The Core Challenge

GPU inference workloads are **memory-bandwidth bound**: time is spent moving data, not computing. Reducing precision (FP32 → FP16) halves traffic—but at what accuracy cost?

## Central research question:

*How does numerical precision format affect GPU inference performance, memory usage, and result accuracy—and can row-wise weight scaling recover accuracy lost to FP16 quantization?*

# Project Goal & Research Questions

**Goal:** Implement a two-layer MLP and benchmark it across **6 precision strategies**, measuring:

- Execution time (ms)
- Throughput (GFLOPS)
- Memory bandwidth (GB/s)
- Memory footprint (MB)
- Accuracy vs. FP32 baseline (MSE)

## Research Questions:

- ① How much faster is FP16 than FP32 on the same custom kernel?
- ② Does row-wise scaling meaningfully reduce FP16 quantization error?
- ③ How much faster is CLBlast GEMM than a hand-written kernel?
- ④ Does CLBlast Mixed (FP16 storage → FP32 compute) offer the best of both worlds?

# Success Criteria

---

<b>Criterion</b>	<b>Definition</b>
Correctness	FP32 output matches expected MLP math
Fair comparison	All 6 modes use the same input data & weights per round
Measurable speedup	CLBlast modes show higher throughput than custom kernels
Accuracy ordering	FP16 + Scale shows lower MSE than plain FP16
Stability	App runs continuously without crashing

---

# FP32 vs. FP16 Precision

## FP32 (Single Precision)

- 32 bits / 4 bytes per number
- $\sim$ 7 decimal digits of precision
- Range:  $\pm 3.4 \times 10^{38}$
- The “gold standard” for inference

## FP16 (Half Precision)

- 16 bits / 2 bytes per number
- $\sim$ 3 decimal digits of precision
- Range:  $\pm 65\,504$
- **Half the memory  $\Rightarrow$  double bandwidth**

### Key Insight

GPU inference is often **memory-bandwidth limited**. Using FP16 moves twice as many values per second through the memory bus  $\Rightarrow$  potential 1.5–2 $\times$  speedup, but with measurable accuracy loss.

# What is GEMM & CLBlast?

## GEMM — General Matrix Multiply:

$$C = \alpha \cdot A \times B + \beta \cdot C$$

- Core operation in every neural network layer
- Two variants: **SGEMM** (FP32), **HGEMM** (FP16)

## CLBlast — OpenCL BLAS library:

- Auto-tuned, tiled, vectorized GEMM kernels
- Selects optimal tile sizes per GPU at first call
- **5–20×** faster than naïve kernels at large sizes

### Why CLBlast is Faster

- ➊ **Tiled shared memory**  
Loads tiles into fast on-chip SRAM
- ➋ **Coalesced access**  
Adjacent threads read adjacent memory
- ➌ **Vectorized loads**  
Reads 4–8 floats per instruction
- ➍ **Auto-tuning**  
Benchmarks many configs, picks the fastest

# The Six Precision Modes

#	Mode	Kernel	Precision	Notes
1	FP32	Custom OpenCL	FP32	Gold standard reference
2	FP16	Custom OpenCL	FP16	~50% memory, some accuracy loss
3	FP16 + Scale	Custom OpenCL	FP16 weights + FP32 scales	Row-wise quantization
4	CLBlast FP32	BLAS SGEMM	FP32	Optimized matrix multiply
5	CLBlast FP16	BLAS HGEMM	FP16	Half-precision BLAS
6	CLBlast Mixed	BLAS SGEMM	FP16 stored → FP32 compute	Best of both worlds

All six modes receive the **same random input and deterministic weights** every round.  
FP32 runs first and stores its output as the **accuracy reference**.

# Mode Comparison at a Glance

Mode	Storage	Compute	Bytes/Elem	Accuracy
FP32	FP32	FP32 (custom)	4	Perfect (reference)
FP16	FP16	FP16 (custom)	2	Some error
FP16 + Scale	FP16 + FP32 scales	FP32 accum.	~2.1	Better than FP16
CLBlast FP32	FP32	FP32 (SGEMM)	4	Perfect (= FP32)
CLBlast FP16	FP16	FP16 (HGEMM)	2	Similar to FP16
CLBlast Mixed	FP16 + FP32 copies	FP32 (SGEMM)	~6	≈ FP32

# Technology Stack

## Backend (Rust)

- `opencl3` — OpenCL bindings
- `libloading` — dynamic DLL loading
- `half` — FP16 ↔ FP32 conversion
- `rand` — random input generation
- `serde` — JSON for Tauri IPC
- `tauri 2.x` — desktop app framework
- `plotters` — PNG chart generation
- `chrono` — timestamps for log folders


## Frontend (TypeScript/React)

- React 19 + Vite 7
- Highcharts — live area charts
- TailwindCSS + DaisyUI
- Tauri IPC via `invoke()`

## GPU

- OpenCL 1.2+ (any GPU vendor)
- CLBlast (embedded DLL)

# System Architecture



# Backend File Structure


File	Responsibility
lib.rs	Tauri commands, CLBlast DLL embedding & loading, device selection, FP16 helpers
types.rs	All structs (MLPInference, RoundData, metrics) + all 6 inference functions
kernel.rs	OpenCL C kernel source code (string constants, compiled at runtime)
logger.rs	CSV export + PNG chart generation (plotters) for automatic logging
main.rs	Minimal entry point — calls lib::run()

## Key design decisions:

- **Singleton GPU context** — avoids repeated OpenCL init (~200 ms each)

# Two-Layer MLP Forward Pass

## Network Architecture



- `batch_size = 64 (fixed)`
- `input_size = hidden_size = matrix_size (128 / 256 / 512 / 1024)`
- `output_size = matrix_size / 2`

**Layer 1:**  $\mathbf{H} = \text{ReLU}(\mathbf{X} \cdot \mathbf{W}_1^T + \mathbf{b}_1)$

**Layer 2:**  $\mathbf{Y} = \mathbf{H} \cdot \mathbf{W}_2^T + \mathbf{b}_2$

Same random  $\mathbf{X}$  and deterministic  $\mathbf{W}_1, \mathbf{W}_2, \mathbf{b}_1, \mathbf{b}_2$  for all 6 modes each round.

# Row-Wise Weight Scaling (Mode 3)

**Problem:** Naïve FP16 conversion loses precision for large weight values.

**Solution:** Normalize each row so all values fit in  $[-1, 1]$  (FP16's sweet spot):

## CPU Pre-processing

$$\text{scale}[h] = \max_i |\mathbf{W}[h][i]| \quad \mathbf{W}_{\text{fp16}}[h] = \text{f32\_to\_f16}\left(\frac{\mathbf{W}[h]}{\text{scale}[h]}\right)$$

## GPU Kernel — recovering original magnitude

$$\text{sum} += (\text{float}) \text{input}[i] \times (\text{float}) \mathbf{W}_{\text{fp16}}[h][i] \times \text{scale}[h]$$

- All FP16 weight values are in  $[-1.0, 1.0]$  — maximum representational density
- Hidden buffer kept as FP32 to avoid a second round of quantization error
- Small overhead: one extra multiply per element + storing the scale vector

# CLBlast GEMM Integration (Modes 4–6)

## CLBlast SGEMM Call (per layer)

```
CLBlastSgemm(RowMajor, NoTrans, Trans, M, N, K, 1.0, A, B, 0.0, C)
```

- **Transpose on  $B$**  (weights) avoids costly memory rearrangement
- BLAS has no concept of bias or ReLU  $\Rightarrow$  separate helper kernels:
  - add\_bias\_relu after Layer 1
  - add\_bias after Layer 2

## CLBlast Mixed mode additionally:

- ① Uploads weights as FP16 (saves bandwidth)
- ② Runs convert\_fp16\_to\_fp32 kernel **on GPU**
- ③ Calls SGEMM in full FP32  $\Rightarrow$  near-FP32 accuracy

## Auto-tuning

First GEMM call per matrix size triggers CLBlast's auto-tuner (1–3 s).

Solved by `warmup_clblast()` before timing begins.

# Performance Metrics

Metric	Unit	Formula / Method
Execution Time	ms	Wall-clock: <code>Instant::now() → queue.finish()</code>
Throughput	GFLOPS	$\frac{\text{batch} \times [H(2I + 1) + O(2H + 1)]}{t \times 10^9}$
Memory BW	GB/s	$\frac{\sum(\text{bytes read} + \text{bytes written})}{t \times 10^9}$
Memory Footprint	MB	$\frac{\text{total GPU buffer bytes}}{1024^2}$
Accuracy (MSE)	—	$\frac{1}{N} \sum_i (\text{output}_i - \text{fp32\_ref}_i)^2$

- FP16 buffers count **2 bytes/ elem**; FP32 buffers count **4 bytes/ elem**
- FP32 and CLBlast FP32 always report MSE = 0.0 (they *are* the reference)
- Timing excludes data generation and DLL loading

# Evaluation Methodology

- ➊ **Shared data:** One RoundData struct (random input + deterministic weights) passed by reference to all 6 modes
- ➋ **Reference first:** FP32 always runs first; its output becomes the accuracy baseline
- ➌ **Warmup:** `warmup_cblast()` fires untimed dummy GEMMs before measurement
- ➍ **Continuous benchmarking:** Self-chaining async loop runs all 6 modes back-to-back, updating 4 live charts per round
- ➎ **Adjustable workload:** Matrix size slider (128–1024) lets you study how problem size affects the tradeoffs

# Automatic Logging & Plotting

Every **5 rounds**, the app automatically saves a full snapshot:

File	Contents
metrics.csv	All metrics from round 1 to current, all 6 modes
execution_time.png	Line chart comparing execution time
throughput.png	GFLOPS comparison
bandwidth.png	Memory bandwidth comparison
accuracy_mse.png	Accuracy MSE vs FP32 baseline

**Log path:** parallel\_log/{YYYY-MM-DD\_HH-MM-SS}\_{matrix\_size}/

- New session folder created automatically when matrix size changes
- Charts generated server-side in Rust via the **plotters** crate
- Colors match the live dashboard (green / blue / purple / yellow / cyan / pink)

# Optimization Techniques

Technique	Where	Effect
Two-pass kernel design	Custom kernels	Each hidden value computed once, not $O$ times
CLBlast auto-tune warmup	<code>warmup_clblast()</code>	Prevents tuning latency in benchmarks
Pre-compiled OpenCL programs	Bias/ReLU + convert	Saves ~50 ms per round
Singleton GPU context	<code>MLP_INSTANCE</code>	Avoids repeated OpenCL init (~200 ms)
<code>spawn_blocking</code>	Tauri handlers	Keeps UI thread responsive
Single React dispatch	<code>chartReducer</code>	One re-render per round, not six
Embedded DLL	<code>include_bytes!</code>	Zero external dependencies

# Challenges Encountered

## CLBlast auto-tuner latency

- First GEMM call: 1–3 s
- Solution: untimed warmup calls

## FP16 in Rust

- No native f16 type
- Stored as u16, converted via half crate
- CLBlast alpha/beta also as u16

## Thread safety of GPU state

- Raw OpenCL handles not Send+Sync
- unsafe impl guarded by mutex

## DLL distribution

- External DLL is fragile
- Solution: include\_bytes! bakes DLL into .exe, extracted to %TEMP% at runtime

# Expected Performance Results

*Approximate relative values at `matrix_size` = 512, `batch_size` = 64. Actual numbers vary by GPU.*

Mode	Memory Footprint	Rel. Throughput	Accuracy MSE
FP32	100% (baseline)	1.0×	0.0 (reference)
FP16	~50%	1.2–1.8×	Small, non-zero
FP16 + Scale	~52%	1.0–1.5×	<b>Less than FP16</b>
CLBlast FP32	~100%	3–10×	0.0
CLBlast FP16	~50%	4–15×	Similar to FP16
CLBlast Mixed	~150%	2–8×	≈ 0.0

CLBlast advantages grow significantly at larger matrix sizes (512+).  
At `matrix_size` = 128, custom kernels may be competitive due to CLBlast overhead.

# Analysis

## Precision Format vs. Speed

FP16 kernels are faster due to **reduced memory bandwidth pressure**—each value is 2 B instead of 4 B. GPU cores run at similar speed for both.

## BLAS vs. Custom Kernels

CLBlast is dramatically faster at large sizes: **tiled shared memory**, vectorized loads, and auto-tuned parameters vs. a simple inner loop.

## Row-Wise Scaling Effect

FP16 + Scale consistently shows **lower MSE** than plain FP16. Most visible when weights have high dynamic range. Small speed cost from extra multiply + FP32 hidden buffer.

## CLBlast Mixed Tradeoff

Uses the most GPU memory (both FP16 + FP32 buffers), but achieves **near-FP32 accuracy**.

# Conclusions

- ① **CLBlast is significantly faster** than hand-written kernels at medium-to-large sizes — specialized BLAS optimizations cannot be matched by a simple loop.
- ② **FP16 halves memory footprint** and produces modest speedups, but accuracy loss is real and measurable.
- ③ **Row-wise scaling meaningfully reduces FP16 error** by keeping values in FP16's most precise range. Accuracy benefit outweighs the small performance cost.
- ④ **CLBlast Mixed offers near-FP32 accuracy at FP16 memory cost** — the pragmatic production choice when accuracy cannot be compromised.
- ⑤ **CLBlast auto-tuning is a one-time cost** that must be accounted for in benchmarks.

# Limitations

- **Windows-only** — CLBlast shipped as .dll; Linux/macOS would need .so/.dylib
- **Fixed batch size (64)** — configurable batch would reveal more tradeoffs
- **No hardware counters** — metrics are analytic, not measured from GPU perf counters
- **Auto-tuning per-session** — CLBlast cache lost on app close
- **Two layers only** — deeper networks may show different error accumulation behavior

## Future Work

- INT8 quantization for further memory reduction
- Deeper networks (4–8 layers) to study error accumulation
- Configurable batch size in the UI
- Cross-platform support (Linux/macOS)
- Hardware counter profiling for true bandwidth measurement
- BF16 (Brain Float 16) as an additional comparison point
- Transformer attention block to study mixed-precision in modern architectures

## Questions Addressed

**Objectives achieved?** Yes — all 6 modes implemented, correct, and benchmarked fairly on the same data per round. Tradeoffs clearly visible in live dashboard.

**Biggest bottleneck?** Custom kernels: **memory bandwidth**. CLBlast vs. custom: lack of **shared-memory tiling**.

**vs. Expectations?** CLBlast delivered **larger speedups** than expected. Row-wise scaling delivered **more accuracy improvement** than expected at almost no throughput cost.

**Do differently?** Use GPU hardware performance counters from the start; expose batch size as a UI parameter.

# Thank You

Mixed-Precision Inference on GPUs

Built with **Rust** · **OpenCL** · **CLBlast** · **Tauri** · **React** · **TypeScript**

GPU Computing Project — Option 2: Mixed Precision with Row-Wise Scaling